

实验六、综合设计实验

实验目的

基于BlobFS，设计一个key-value数据库，支持open，put，get，close操作

实验内容

基于BlobFS，设计一个key-value数据库，支持open，put，get，close操作

实验过程和步骤

实验思路

本次基于BlobFS设计KV数据库。考虑到BlobFS仅支持追加写的特性，本次实验选择使用C++复现Bitcask数据库。

- BlobFS的限制

Limitations

- BlobFS has primarily been tested with RocksDB so far, so any use cases different from how RocksDB uses a filesystem may run into issues. BlobFS will be tested in a broader range of use cases after this initial release.
- Only a synchronous API is currently supported. An asynchronous API has been developed but not thoroughly tested yet so is not part of the public interface yet. This will be added in a future release.
- File renames are not atomic. This will be fixed in a future release.
- BlobFS currently supports only a flat namespace for files with no directory support. Filenames are currently stored as xattrs in each blob. This means that filename lookup is an $O(n)$ operation. An SPDK btree implementation is underway which will be the underpinning for BlobFS directory support in a future release.
- Writes to a file must always append to the end of the file. Support for writes to any location within the file will be added in a future release.

Bitcask数据库原理

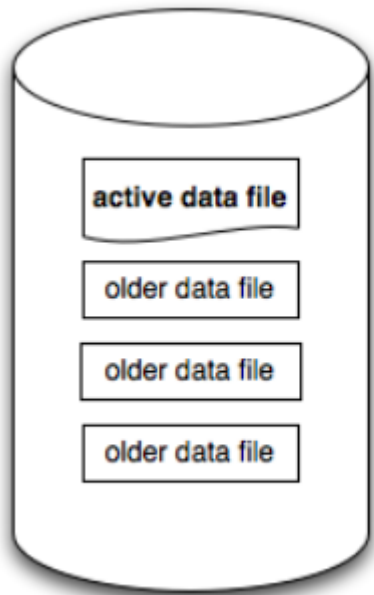
日志型数据存储

所有写操作只追加而不修改老的数据，这样做目的是能保证最大程度的顺序 IO，压榨出机械硬盘的顺序写性能。

在Bitcask模型中，数据文件以日志型只增不减的写入文件，而文件有一定的大小限制，当文件大小增加到相应的限制时，就会产生一个新的文件，老的文件将只读不写。

在任意时间点，只有一个文件是可写的，在Bitcask模型中称其为active data file，而其他的已经达到限制大小的文件，称为older data file，如下图：

a bitcask on disk



文件中的数据结构非常简单，是一条一条的数据写入操作，每一条数据的结构如下：


crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value
crc	tstamp	ksz	value_sz	key	value

基于hash表的索引数据

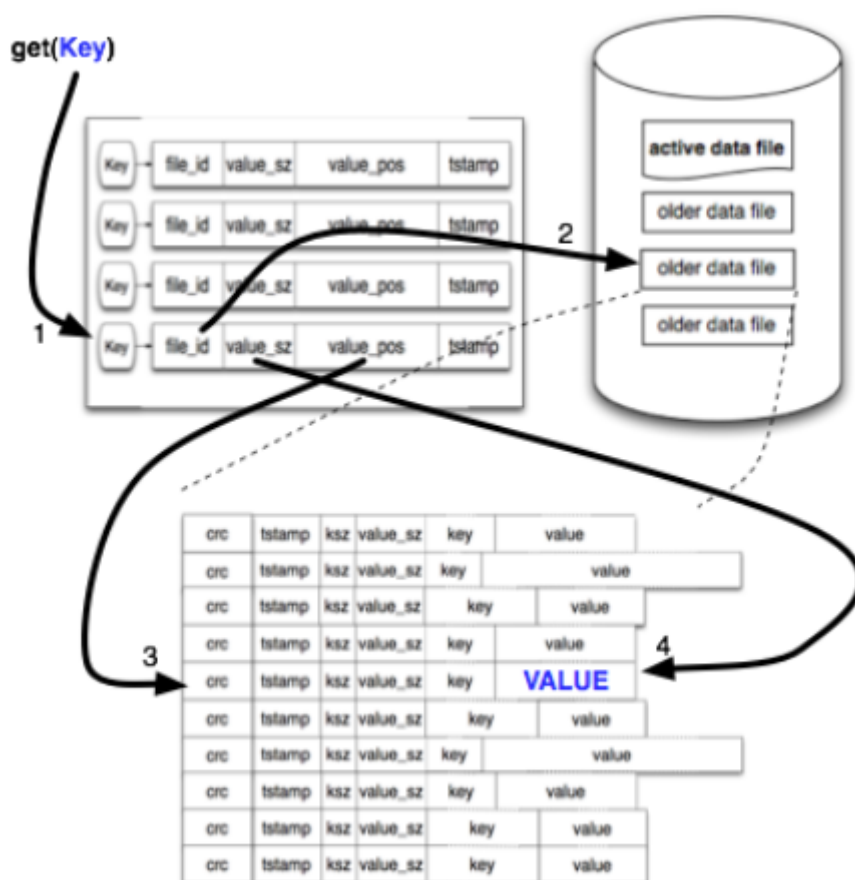
日志类型的数据文件会让我们的写入操作非常快，而如果在这样的日志型数据上进行key值查找，那将是一件非常低效的事情。于是我们需要使用一些方法来提高查找效率。

例如在Bigtable中，使用bloom-filter算法为每一个数据文件维护一个bloom-filter 的数据块，以此来判定一个值是否在某一个数据文件中。

在Bitcask模型中，除了存储在磁盘上的数据文件，还有另外一块数据，那就是存储在内存中的hash表，hash表的作用是通过key值快速的定位到value的位置。hash表的结构大致如下图所示：



hash表对应的这个结构中包括了三个用于定位数据value的信息，分别是文件id号(file_id)，value值在文件中的位置(value_pos)，value值的大小(value_sz)，于是我们通过读取file_id对应文件的value_pos开始的value_sz个字节，就得到了我们需要的value值。整个过程如下图所示：



由于多了一个hash表的存在，我们的写操作就需要多更新一块内容，即这个hash表的对应关系。于是一个写操作就需要进行一次顺序的磁盘写入和一次内存操作。

代码实现

bitcask.h

```
#ifndef BITCASK_H_
#define BITCASK_H_

#include <iostream>
#include <vector>
#include <string>
#include <unordered_map>
#include <pwd.h>
#include <boost/archive/binary_oarchive.hpp>
#include <boost/archive/binary_iarchive.hpp>
#include <boost/serialization/string.hpp>
#include <boost/serialization/vector.hpp>
#include <boost/serialization/list.hpp>
#include <boost/crc.hpp>
#include <boost/date_time/posix_time/posix_time.hpp>
#include <boost/date_time/posix_time/time_serialize.hpp>

using namespace std;

#define filemax 1024 * 4096 // 4MB

const string fileprev = "bitcask_data";
const string cmd_prompt = ">>> bitcask : ";
const string cmd = ">>> ";
const int number = 0;
// data
struct bitcask_data
{
    string key;
    int key_len;
    string value;
    int value_len;
    boost::posix_time::ptime timestamp;
    uint32_t crc;
    // crc
    template <typename Archive>
    void serialize(Archive &ar, const unsigned int version)
    {
        ar &key_len;
        ar &key;
        ar &value_len;
        ar &value;
        ar &crc;
        ar &timestamp;
    }
};

// index
struct bitcask_index
```

```

{
    string key;
    string file_id;
    int value_pos;
    int value_len;
    boost::posix_time::ptime timestamp;
    bool value_valid;

    template <typename Archive>
    void serialize(Archive &ar, const unsigned int version)
    {
        ar &key;
        ar &file_id;
        ar &value_pos;
        ar &value_len;
        ar &timestamp;
        ar &value_valid;
    }

    bitcask_index()
    {
        value_valid = false;
    }
};

// bitcask
class bitcask
{
private:
    unordered_map<string, bitcask_index> index;
    int _activefile;
    bool _start;
    bool _finish;
    string _response;
    string filepath;

private:
    void init(string path);
    uint32_t crc32(string value);
    void insert_data(string key, string value);
    void write_data(bitcask_data newdata);
    void write_index(bitcask_index newindex);
    bitcask_data read_data(string key);
    void read_datainfo(string key);
    bitcask_index read_index(string key);
    void delete_data(string key);
    void update_data(string key, string value);
    void update_index(bitcask_index upindex, string key);
    void merge();
    void flush(); // flush index :hint.bin

public:
    bitcask();
    string Get(string key);

```

```
void Put(string key, string value);
void Open(string path);
void Close();
~bitcask();

};

#endif /* BITCASK_H_ */
```

bitcask.cpp

```
#include <iostream>
#include <fstream>
#include <string>
#include <ctime>
#include <unistd.h>
#include <dirent.h>
#include <pwd.h>
#include "bitcask.h"
using namespace std;

bitcask::bitcask()
{
    _start = false;
    _activefile = 0;
    _finish = false;
    _response = "";
    filepath = "";
}

bitcask::~bitcask()
{
    if (this->_start)
    {
        merge();
        flush();
    }
}

void bitcask::init(string path)
{
    this->filepath = path;
    this->_start = true;
    long len;
    fstream hint;
    hint.open(filepath + "hint.bin", ios::binary | ios::out | ios::app);
    if (!hint)
    {
```

```

        cout << "the file hint.bin open failure or maybe not exist!\n";
    }
    len = hint.tellg();
    if (len == 0)
    {
        cout << "create file hint.bin successful!\n";
    }
    else
    {
        /*
         * load index to memory
         */
        // cout << "loading index" << endl;
        bitcask_index search;
        fstream hint;
        hint.open(filepath + "hint.bin", ios::binary | ios::in);
        if (!hint)
        {
            cout << "the file hint.bin open failure or maybe not exist!\n";
        }

        while (hint)
        {
            boost::archive::binary_iarchive ia(hint,
boost::archive::no_header);
            try
            {
                ia >> search;
            }
            catch (const exception &e)
            {
                // cout << "read end" << endl;
                goto do_load;
            }
            // cout << "loading:" << search.key << endl;
            bitcask_index insert;
            insert.file_id = search.file_id;
            insert.value_pos = search.value_pos;
            insert.value_len = search.value_len;
            insert.timestamp = search.timestamp;
            insert.value_valid = search.value_valid;
            index[search.key] = insert;
        }
    }
do_load:
    fstream filelog;
    filelog.open(filepath + "filelog.bin", ios::binary | ios::in);
    if (!filelog)
    {
        cout << "the file filelog.bin open failure or maybe not exist!\n";
    }
    filelog.read((char *)&_activefile, sizeof(int));
    filelog.close();
    if (_activefile == 0)

```

```

{
    cout << "create file filelog.bin successful!\n";
    _activefile = 1;
    filelog.open(filepath + "filelog.bin", ios::binary | ios::out |
ios::app);
    filelog.write((char *)&_activefile, sizeof(int));
    filelog.close();
    return;
}
_start = true;
}

uint32_t bitcask::crc32(string value)
{
    boost::crc_32_type result;
    result.process_bytes(value.c_str(), value.length());
    return result.checksum();
}

void bitcask::insert_data(string key, string value)
{
    bitcask_index search = read_index(key);
    if (search.key != "")
    {
        cout << "the data " + key + " already exist!\n";
        return update_data(key, value);
    }
    // add data
    bitcask_data newdata;
    newdata.key = key;
    newdata.key_len = int(key.length());
    newdata.value = value;
    newdata.value_len = int(value.length());
    newdata.crc = crc32(value);
    // newdata.timestamp=time(0);
    newdata.timestamp = boost::posix_time::microsec_clock::universal_time();

    // add index
    fstream datafile;
    bitcask_index newindex;
    newindex.key = key;
    newindex.file_id = fileprev + to_string(_activefile);
    datafile.open(filepath + newindex.file_id, ios::binary | ios::out |
ios::app);
    if (!datafile)
        cout << cmd_prompt + newindex.file_id + " open failure\n";
    newindex.value_pos = datafile.tellg();
    if (newindex.value_pos > filemax || filemax - newindex.value_pos <
sizeof(newdata))
    {
        _activefile++;
        newindex.file_id = fileprev + to_string(_activefile);
    }
    newindex.timestamp = newdata.timestamp;
}

```



```

        newindex.value_len = sizeof(newdata);
        newindex.value_valid = true;
        datafile.close();
        write_data(newdata);
        write_index(newindex);
        // add to memory index array
        index[key] = newindex;
        // cout << "the data " + key + " insert successful\n";
    }

void bitcask::write_data(bitcask_data newdata)
{
    string file = fileprev + to_string(_activefile);
    fstream datafile;
    datafile.open(filepath + file, ios::binary | ios::out | ios::app);
    if (!datafile)
        cout << file + " open file " + file + " failure!\n";
    // datafile.write((char *)&newdata, sizeof(newdata));
    // data_append(datafile, newdata);
    boost::archive::binary_oarchive oa(datafile, boost::archive::no_header);
    oa << newdata;
    datafile.close();
}

void bitcask::write_index(bitcask_index newindex)
{
    fstream hint;
    hint.open(filepath + "hint.bin", ios::binary | ios::out | ios::app);
    if (!hint)
    {
        cout << "the file hint.bin open failure!\n";
    }
    // hint.write((char *)&newindex, sizeof(newindex));
    // cout << "writing index: " << newindex.key << newindex.timestamp <<
endl;
    boost::archive::binary_oarchive oa(hint, boost::archive::no_header);
    oa << newindex;
    hint.close();
}

bitcask_index bitcask::read_index(string key)
{
    bitcask_index search;
    if (" " == key)
    {
        return search;
    }
    for (auto indexinfo : index)
    {
        if (indexinfo.first == key)
        {
            {

```

```

        search.key = indexinfo.first;
        search.file_id = indexinfo.second.file_id;
        search.value_pos = indexinfo.second.value_pos;
        search.value_len = indexinfo.second.value_len;
        search.value_valid = indexinfo.second.value_valid;
        search.timestamp = indexinfo.second.timestamp;
        // cout << "got key for " << key << " ,valid " <<
search.value_valid << endl;
        return search;
    }
}
return search;
}

bitcask_data bitcask::read_data(string key)
{
    bitcask_data search_data;
    bitcask_index search_index = read_index(key);
    if (search_index.value_valid == true)
    {
        string file = search_index.file_id;
        // cout << "reading file" << file << endl;
        fstream datafile;
        datafile.open(filepath + file, ios::binary | ios::in);
        if (!datafile)
            cout << "open file " + file + " failure\n";

        boost::archive::binary_iarchive ia(datafile,
boost::archive::no_header);
        if (search_index.value_pos)
        {
            datafile.seekg(search_index.value_pos, ios::beg);
            // cout << "seeking to" << search_index.value_pos << endl;
        }
        // datafile.read((char *)&search_data, sizeof(search_data));
        // cout << "reading datafile" << filepath + file << endl;
        // cout << "key = " << search_index.key << endl;
        // cout << "value_pos = " << search_index.value_pos << endl;
        // cout << "ts = " << search_index.timestamp << endl;
        // cout << "file pos = " << datafile.tellp() << endl;
        // search_data.serialize(ia, 0);
        ia >> search_data;
        // cout << "file pos = " << datafile.tellp() << endl;
        // cout << "on-disk ts = " << search_data.timestamp << endl;
        // return search_data;
    }
    // else
    //     cout << "the data " + key + " does not exist!\n";
    return search_data;
}

void bitcask::read_datainfo(string key)
{

```

```

        bitcask_data data = read_data(key);
        bitcask_index index = read_index(key);
        if (index.value_valid == true)
        {
            _response += cmd + "key :" + key + "\n";
            _response += cmd + "value :" + data.value + "\n";
            _response += cmd + "crc :" + to_string(data.crc) + "\n";
            _response += cmd + "file id :" + index.file_id + "\n";
            _response += cmd + "value pos :" + to_string(index.value_pos) + "\n";
            _response += cmd + "value length :" + to_string(data.value_len) +
"\n";

            // cout << _response;
            //_response+=cmd+"time :"+ data.timestamp;
            //_response+=cmd+"time :"+ctime(&data.timestamp);
        }
        else
            return;
    }

void bitcask::delete_data(string key)
{
    bitcask_index delindex = read_index(key);
    if (delindex.key != "")
    {
        delindex.value_valid = false;
        index[key] = delindex;
        cout << "the data " + key + " already delete!\n";
    }
    else
        _response += cmd_prompt + "the data " + key + " does not exist!\n";
}

void bitcask::update_data(string key, string value)
{
    bitcask_index search = read_index(key);
    if (search.key == "")
    {
        cout << "the data " + key + " does not exist!\n";
        return;
    }

    // update data
    fstream datafile, hintfile;
    bitcask_data updata;
    bitcask_index upindex = read_index(key);
    updata.key = key;
    updata.key_len = int(key.length());
    updata.value = value;
    updata.value_len = int(value.length());
    // updata.timestamp=time(0);
    updata.timestamp = boost::posix_time::microsec_clock::universal_time();

    // update index
    upindex.file_id = fileprev + to_string(_activefile);

```

```

        datafile.open(filepath + upindex.file_id, ios::binary | ios::in |
ios::app);
        if (!datafile)
            _response += cmd_prompt + "the file " + upindex.file_id + " open
failure!\n";
        upindex.value_pos = datafile.tellg();
        if (upindex.value_pos > filemax || filemax - upindex.value_pos <
sizeof(updata))
        {
            _activefile++;
        }
        upindex.timestamp = updata.timestamp;
        upindex.value_len = sizeof(updata);
        upindex.value_valid = true;
        datafile.close();
        write_data(updata);
        update_index(upindex, key);
        _response += cmd_prompt + "the data " + key + " update successful\n";
    }

void bitcask::update_index(bitcask_index upindex, string key)
{
    bitcask_index seaindex = read_index(key);
    seaindex.key = key;
    seaindex.file_id = upindex.file_id;
    seaindex.value_pos = upindex.value_pos;
    seaindex.value_len = upindex.value_len;
    seaindex.value_valid = upindex.value_valid;
    seaindex.timestamp = upindex.timestamp;
    index[key] = seaindex;
}

void bitcask::merge()
{
    /*
    merge data in file
    function: delete data in file
    */
    int beans = 1;
    long value_pos;
    vector<bitcask_data> data_array;
    // cout << "merge begin: activefile" << _activefile << endl;
    for (; beans ≤ _activefile; beans++)
    {
        string file = fileprev + to_string(beans);
        fstream datafile;
        datafile.open(filepath + file, ios::binary | ios::in);

        if (!datafile)
        {
            _response += cmd_prompt + "the data file " + file + " open
failure!\n";
        }
    }
}

```

```

        bitcask_data beans_data;
        bitcask_index beans_index;
        datafile.seekg(0, ios::beg);
        // load to memory
        while (datafile)
        {
            boost::archive::binary_iarchive ia(datafile,
boost::archive::no_header);
            try
            {
                ia >> beans_data;

                beans_index = read_index(beans_data.key);
                // cout << "ts:" << beans_data.timestamp << " vs " <<
beans_index.timestamp << endl;
            }
            catch (const exception &e)
            {
                goto do_merge;
            }

            if (beans_index.value_valid == true && beans_data.timestamp ==
beans_index.timestamp)
            {
                // cout << "pushing:" << beans_index.key << endl;
                data_array.push_back(beans_data);
            }
        }

do_merge:
    for (auto data : data_array)
    {
        // cout << "dataarray = " << data.key << endl;
    }

    datafile.close();
    // write to file
    datafile.open(filepath + file, ios::binary | ios::out);
    if (!datafile)
    {
        _response += cmd_prompt + "the data file " + file + " open
failure!\n";
    }
    datafile.seekg(0, ios::beg);

    for (auto data : data_array)
    {
        value_pos = datafile.tellg();
        // datafile.write((char *)(&data), sizeof(data));
        boost::archive::binary_oarchive oa(datafile,
boost::archive::no_header);
        oa << data;
        bitcask_index seaindex = read_index(data.key);
        seaindex.value_pos = value_pos;
    }

```

```

        index[data.key] = seaindex;
    }
    datafile.close();
    data_array.clear();
}
/*
  TODO merge file
*/
if (_activefile ≥ 2)
{
    while (_activefile > 1)
    {
        string file = fileprev + to_string(_activefile);
        fstream datafile;
        datafile.open(filepath + file, ios::binary | ios::in);
        if (!datafile)
        {
            _response += cmd_prompt + "the data file " + file + " open
failure!\n";
        }
        bitcask_data beans_data;
        bitcask_index beans_index;
        datafile.seekg(0, ios::beg);
        // TODO load to memory maybe too big?
        while (datafile)
        {
            boost::archive::binary_iarchive ia(datafile,
boost::archive::no_header);
            try
            {
                ia >> beans_data;
            }
            catch (const exception &e)
            {
                goto do_merge_2;
            }

            data_array.push_back(beans_data);
        }

        do_merge_2:
        datafile.close();
        // write to file
        for (int pos = 1; pos < _activefile; pos++)
        {
            string mergefile = fileprev + to_string(pos);
            fstream datafile;
            datafile.open(filepath + mergefile, ios::binary | ios::out |
ios::app);
            if (!datafile)
            {
                _response += cmd_prompt + "the data file " + file + "
open failure!\n";
            }

```

```

        long mergefile_end = datafile.tellg();
        bitcask_data merge_data = data_array.back();
        bitcask_index merge_index = read_index(merge_data.key);
        while (mergefile_end < filemax && (filemax - mergefile_end) <
sizeof(merge_data))
        {
            // datafile.write((char *)&merge_data,
sizeof(merge_data));
            boost::archive::binary_oarchive oa(datafile,
boost::archive::no_header);
            oa << merge_data;
            merge_data = data_array.back();
            merge_index = read_index(merge_data.key);
            merge_index.value_pos = mergefile_end;
            merge_index.file_id = mergefile;
            index[merge_data.key] = merge_index;
            mergefile_end += sizeof(merge_data);
            data_array.pop_back();
            merge_data = data_array.back();
        }
        datafile.close();
    }
    if (data_array.size() == 0)
    {
        _activefile--;
    }
    else
    {
        fstream newdatafile;
        newdatafile.open(filepath + file, ios::binary | ios::out);
        if (!newdatafile)
        {
            _response += cmd_prompt + "the data file " + file + "
open failure!\n";
        }
        newdatafile.seekg(0, ios::beg);
        for (auto data : data_array)
        {
            // newdatafile.write((char *)&data, sizeof(data));
            boost::archive::binary_oarchive oa(newdatafile,
boost::archive::no_header);
            oa << data;
        }
        newdatafile.close();
        break;
    }
}
}

void bitcask::flush()
{
    // write index file to index file hint.bin
    fstream hint;

```

```

hint.open(filepath + "hint.bin", ios::binary | ios::out);
if (!hint)
{
    _response += cmd_prompt + "the file hint.bin open failure!\n";
}
hint.seekg(0, ios::beg);
for (auto indexinfo : index)
{
    if (indexinfo.second.value_valid == true)
    {
        // hint.write((char *)&indexinfo.second,
sizeof(indexinfo.second));
        boost::archive::binary_oarchive oa(hint,
boost::archive::no_header);
        oa << indexinfo.second;
    }
}
// hint.close();
// write active file number to file filelog.bin
fstream filelog;
filelog.open(filepath + "filelog.bin", ios::binary | ios::out);
if (!filelog)
{
    _response += cmd_prompt + "the filelog.bin open failuer!\n";
}
filelog.write((char *)&_activefile, sizeof(int));
}

string bitcask::Get(string key)
{
    if (this->_start == true)
    {
        bitcask_data bc_data = read_data(key);
        if (bc_data.key.length() == 0)
            cout << key + " does not exist!" << endl;
        return bc_data.value;
    }
    else
    {
        cout << "please open a database first" << endl;
        return "";
    }
}

void bitcask::Put(string key, string value)
{
    if (this->_start == true)
        insert_data(key, value);
    else
        cout << "please open a database first" << endl;
}

void bitcask::Open(string path)
{

```



```

        if (this->_start == false)
            init(path);
        else
            cout << "already open a database, please close first" << endl;
    }

    void bitcask::Close()
    {
        if (this->_start)
        {
            // merge();
            flush();
            _start = false;
            _activefile = 0;
            _finish = false;
            _response = "";
            filepath = "";
        }
        else
            cout << "please open a database first" << endl;
    }

```

main.cpp

```

#include "bitcask.h"

#define db_path "/home/miracle/work/task6/bitcask/db/"

int main()
{
    bitcask *db = new bitcask;
    db->Open(db_path);

    db->Put("key1", "value1");
    cout << db->Get("key1") << endl;
    db->Close();
    cout << db->Get("key1") << endl;
    db->Open(db_path);
    cout << db->Get("key1") << endl;
    cout << db->Get("key2") << endl;
    return 0;
}

```

实验结果

挂载BlobFS

将BlobFS挂载到当前目录下的db文件夹，用作数据库存储

```
root@cs-exp-zns:/home/miracle/work/spdk/test/blobfs# ./fuse/fuse ./nvme.json Nvme0n1 /home/miracle/work/task6/bitcask/db/
[2022-12-22 20:09:33.675484] Starting SPDK v23.01-pre / DPOK 22.07.0 initialization...
[2022-12-22 20:09:33.675597] [ DPOK EAL parameters: [2022-12-22 20:09:33.675611] spdk fuse [2022-12-22 20:09:33.675621] --no-shconf [2022-12-22 20:09:33.675628] -c 0x3
[2022-12-22 20:09:33.675636] --huge-unlink [2022-12-22 20:09:33.675644] --log-level=lib.eal:6 [2022-12-22 20:09:33.675653] --log-level=lib.cryptodev:5 [2022-12-22 20:09
:33.675678] --log-level=user1:6 [2022-12-22 20:09:33.675687] --iova-mode=pa [2022-12-22 20:09:33.675696] --base-virtaddr=0x200000000000 [2022-12-22 20:09:33.675704] --m
atch-allocations [2022-12-22 20:09:33.675713] --file-prefix=spdk_pid1149 [2022-12-22 20:09:33.675721] ]
TELEMETRY: No legacy callbacks, legacy socket not created
[2022-12-22 20:09:33.803886] app.c: 705:spdk app start: *NOTICE*: Total cores available: 2
[2022-12-22 20:09:33.832280] reactor.c: 926:reactor_run: *NOTICE*: Reactor started on core 1
[2022-12-22 20:09:33.832280] reactor.c: 926:reactor_run: *NOTICE*: Reactor started on core 0
[2022-12-22 20:09:33.834657] accel_sw.c: 466:sw accel module init: *NOTICE*: Accel framework software module initialized.
Mounting filesystem on bdev Nvme0n1 to path /home/miracle/work/task6/bitcask/db/...
done.
[2022-12-22 20:09:33.985482] blobfs_fuse.c: 239:fuse_loop_new_thread: *NOTICE*: Start to loop blobfs on bdev Nvme0n1 mounted at /home/miracle/work/task6/bitcask/db/
```

运行测试程序

```
root@cs-exp-zns:/home/miracle/work/task6/bitcask# make run
create file hint.bin successful!
the file filelog.bin open failure or maybe not exist!
create file filelog.bin successful!
value1
please open a database first

value1
key2 does not exist!
```

```
int main()
{
    bitcask *db = new bitcask;
    db->Open(db_path);

    db->Put("key1", "value1");
    cout << db->Get("key1") << endl;
    db->Close();
    cout << db->Get("key1") << endl;
    db->Open(db_path);
    cout << db->Get("key1") << endl;
    cout << db->Get("key2") << endl;
    return 0;
}
```

对比测试流程，可以看到

1. 先打开了数据库，程序检测到该处没有任何文件于是进行初始化，并提示hint及filelog文件创建成功
2. `Put("Key1", "value1")`
3. `Read("Key1")` 此时输出了对应的值 `value1`
4. 关闭数据库
5. 再次读取 `Key1` 的值（`Read("Key1")`），提示数据库未打开
6. 重新打开数据库
7. 分别获取 `Key1`，`Key2` 的值
8. 可以看到获取 `Key1` 值时输出了正确的 `value1`，获取 `Key2` 值时提示 `Key2` 不存在

经过以上测试通过了Open,Close,Put,Get等基本语法要求

检查二进制存储文件

```
root@cs-exp-zns:/home/miracle/work/task6/bitcask# hexdump -C ./db/bitcask_data1
00000000  00 00 00 00 00 04 00 00 00 04 00 00 00 00 00 00 |.....|
00000010  00 6b 65 79 31 06 00 00 00 06 00 00 00 00 00 00 |.key1.....|
00000020  00 76 61 6c 75 65 31 5a 6c 75 a2 00 00 00 00 00 |.value1Zlu.....|
00000030  00 00 00 00 00 08 00 00 00 00 00 00 00 32 30 32 |.....202|
00000040  32 31 32 32 32 00 01 00 00 00 00 0c 00 00 00 00 |21222.....|
00000050  00 00 00 0b 00 00 00 00 00 00 00 05 00 00 00 00 |.....|
00000060  00 00 00 e4 0e 06 00 00 00 00 00 |.....|
0000006b
```

可以看出符合之前设计的存储结构

实验结论和心得体会

本次实验学习了BlobFS相关特性并在其上实现了KV数据库的Open,Close,Put,Get操作。